

DAY 3: FEATURE ENGINEERING & STATISTICAL
MODEL EVALUATION

STATUS: Fully Executed

DATASET: Adult Census (N=32,561)

TARGET: Income >50K (24.08%)

A principled approach to expanding feature representation without data leakage, cross-validating competing algorithms, and benchmarking dimensionality trade-offs.

1. Task 1: Engineered Features Dictionary & Univariate Signals

Eight domain-principled features were constructed strictly at row-level (zero target leakage). Mutual Information (MI) and grouped target rates confirm substantial non-linear predictive signals:

Feature Name	Type	Creation Rule (Zero-Leakage)	Domain Justification / Signal Logic	MI	Univariate Signal (P(>50K))
is_married	Binary Flag	<code>marital_status in {Civ-Spouse, AF-Spouse}</code>	Household stability & dual-income earner proxy; strongest adult census predictor.	0.1105	Unmarried: 6.5% Married: 44.7% (6.9x lift)
edu_x_hours	Interaction	<code>education_num * hours_per_week</code>	Captures compounding marginal wage returns per hour for educated workers.	0.0840	Q1: 6.2% Q2: 16.4% Q3: 30.0% Q4: 52.2%
log_capital_gain	Log Numeric	<code>np.log1p(capital_gain)</code>	Compresses severe right-skewed tail, linearizing massive positive returns.	0.0824	Zero: 20.7% Mid: 89.9% High: 98.8%
age_bucket	Categorical	<code>pd.cut(age, [0, 24, 34, 44, 54, 64, 200])</code>	Models non-linear human earning curve peaking in mid-to-late career stages.	0.0631	<25: 1.1% 25-34: 16.8% 45-54: 40.1% 65+: 20.7%
is_higher_ed	Binary Flag	<code>education in {Bachelors, Masters, Prof, Doc}</code>	Isolates step-function wage credential premium over standard secondary schooling.	0.0494	Non-Higher: 16.1% Higher Ed: 48.5% (3.0x lift)
hours_bin	Categorical	<code>pd.cut(hours, [-1, 20, 35, 40, 50, 200])</code>	Distinguishes non-linear part-time, standard 40h, and extreme overtime pay tiers.	0.0385	Part-time: 6.7% Overtime: 39.6% Extreme (>50h): 41.3%
has_capital_gain	Binary Flag	<code>(capital_gain > 0).astype(int)</code>	Separates asset holders from non-investors given 92% structural sparsity.	0.0323	No Gain: 20.7% Has Gain: 61.8% (3.0x lift)
has_capital_loss	Binary Flag	<code>(capital_loss > 0).astype(int)</code>	Indicator of active portfolio management & liquid wealth exposure.	0.0098	No Loss: 22.8% Has Loss: 50.9% (2.2x lift)

2. Task 2: Rebuilding the Preprocessing Pipeline (Strict Zero-Leakage)

Architecture: Engineered features are generated inside a clean scikit-learn transformer (FeatureEngineer / FunctionTransformer) placed upstream of a ColumnTransformer. Numerical features (including edu_x_hours and log_capital_gain) receive SimpleImputer(strategy='median') and StandardScaler. Categorical features (including age_bucket and hours_bin) receive SimpleImputer(strategy='most_frequent') and OneHotEncoder(handle_unknown='ignore', sparse_output=False). All estimators are fitted strictly on training folds inside cross-validation, guaranteeing zero test leakage.

3. Task 3: 5-Fold Stratified Cross-Validation Benchmark

Three diverse classifiers were evaluated under identical 5-fold stratified CV splits (80/20 train/validation ratio per fold). HistGradientBoosting clearly outperforms competing architectures across all evaluation dimensions:

Model Architecture	Accuracy (Mean ± Std)	ROC-AUC (Mean ± Std)	F1-Score (Mean ± Std)	Primary Metric Assessment & Behavior
HistGradientBoosting	0.8723 ± 0.0032	0.9272 ± 0.0016	0.7120 ± 0.0070	Top Performer. Lowest variance across folds; handles mixed non-linear interactions natively.
Random Forest (100 trees)	0.8557 ± 0.0041	0.9042 ± 0.0032	0.6775 ± 0.0090	Strong recall but lower ROC-AUC precision; higher fold-to-fold variance (±0.0032).
Logistic Regression (L2)	0.8512 ± 0.0028	0.9055 ± 0.0022	0.6589 ± 0.0073	Fast and stable baseline; limited by linear hyperplanes on complex non-linear combinations.





Figure 1: Stratified 5-Fold CV Score Distributions (Accuracy, ROC-AUC, F1). HistGradientBoosting demonstrates clear dominance with zero fold overlap in ROC-AUC.

4. Task 4: Statistical Hypothesis Testing & Feature Importance

To confirm whether HistGradientBoosting's superiority is statistically believable and practically meaningful, paired tests were executed across identical 10-fold CV splits against Logistic Regression:

Hypothesis Test Metric	Value / Result	Statistical & Practical Interpretation
HGB vs LR Mean ROC-AUC	0.9280 vs 0.9143	Mean ROC-AUC Lift: +0.0138 (+1.38%) . Consistent lift across every partition.
Paired t-Test (Parametric)	t = 33.45, p = 9.41e-11	p << 0.001. Null hypothesis (equal means) definitively rejected.
Wilcoxon Signed-Rank	W = 0.0, p = 0.00195	Non-parametric test confirms zero rank inversions: 10 Wins, 0 Losses, 0 Ties .
Practical Significance	Meaningful Lift	In income classification, a +1.38% AUC and +5.3% F1 lift translates to hundreds of additional high-value earners correctly classified without false-positive inflation.



Engineered Feature	HGB Perm. Imp.	LR Coeff (Odds Ratio)	Domain & Mechanical Role
edu_x_hours	0.0237 (Top Eng)	-0.283 (OR: 0.75)*	Highest HGB engineered signal; captures compounding hourly education returns.
log_capital_gain	0.0000 (Subsumed)	+5.170 (OR: 176.0)	Dominant linear driver for LR; massive positive odds for capital market investors.
has_capital_gain	0.0000 (Subsumed)	-5.003 (OR: 0.007)	Acts as negative intercept calibration for non-investors in logistic equation.
age_bucket	0.0012	<25: -1.05 45-54: +0.55	Reflects peak earning years (45-54) vs early entry-level suppression (<25).
hours_bin	0.0003	Reduced: -1.00 Extr: +0.61	Strong threshold demarcation between part-time and excessive overtime work.
is_married	0.0001	+0.437 (OR: 1.55)	HGB splits directly on raw marital_status (imp: 0.0910); LR relies on binary.

Figure 2: 10-Fold Paired ROC-AUC Boxplot (Left). Feature Importance Comparison Table (Right). *Note: LR coefficient for edu_x_hours reflects collinearity adjustment in the presence of base education_num and hours_per_week.

5. Task 5: Feature Selection Benchmark & Dimensionality Check

L1-regularized feature selection (SelectFromModel with LogisticRegression penalty='l1') was benchmarked against the full feature representation across 5-fold CV on HistGradientBoosting:

Feature Set Configuration	Orig / Transformed	Retained Features	Accuracy (Mean ± Std)	ROC-AUC (Mean ± Std)	F1-Score (Mean ± Std)	CV Wall Time
Full Feature Set (All)	14 / 123	123 (100%)	0.8742 ± 0.0036	0.9283 ± 0.0016	0.7154 ± 0.0080	7.92s
Moderate L1 Selection (C=0.1)	14 / 123	49 (39.8%)	0.8739 ± 0.0037	0.9283 ± 0.0015	0.7143 ± 0.0084	8.84s (+0.92s)
Aggressive L1 Selection (C=0.01)	14 / 123	18 (14.6%)	0.8686 ± 0.0022	0.9249 ± 0.0023	0.6983 ± 0.0061	2.88s (-5.04s)

6. Architectural Decision & Hyperparameter Tuning Roadmap (Day 4 Plan)

FINAL ARCHITECTURAL RECOMMENDATION: Selected Algorithm: HistGradientBoostingClassifier. Highly statistically superior to Logistic Regression (p=9.4e-11, 10/10 fold wins) and Random Forest (+2.3% ROC-AUC). Offers native binning, fast histogram computation, and automatic non-linear split handling. Selected Feature Representation: Full Feature Set (All 123 features). Preserves peak ROC-AUC (0.9283) and F1 (0.7154). Moderate L1 selection eliminated 74 features with equal ROC-AUC but increased total runtime (8.84s vs 7.92s) due to the selector fitting step. Aggressive L1 sacrificed 0.34% ROC-AUC and 1.71% F1 score. Because tree boosting inherently performs greedy coordinate feature selection at split points, the full representation maximizes interaction discovery with zero speed penalty.	DAY 4 HYPERPARAMETER SEARCH GRID: - learning_rate: [0.03, 0.05, 0.10, 0.15] - max_leaf_nodes: [15, 31, 63, 127] - min_samples_leaf: [20, 50, 100] - l2_regularization: [0.0, 0.5, 1.0, 5.0] - max_iter: Early stopping with tolerance=1e-4 Objective: Maximize cross-validated ROC-AUC while fine-tuning threshold calibration for F1 optimization.
--	---