

Week 2 Day 4 — CrewAI: Multi-Agent Collaboration & Delegation

Author: Armish Iqbal | System: Enterprise SaaS Competitive Intelligence & TCO Modeling

Executive Summary: Transitioning from single-agent stateful graph architectures (LangGraph) to role-based multi-agent collaboration (CrewAI). We deploy a specialized crew of three autonomous agents (Researcher, Financial Analyst, and Marketing Strategist) with strictly confined tool privileges to solve an enterprise business workflow: auditing SaaS competitor data, computing deterministic multi-team TCO models, and synthesizing an executive sales battlecard. Both **Process.sequential** and **Process.hierarchical** are evaluated across quality, latency, token consumption, and cost.

1. Multi-Agent Design Thinking & Persona Construction

Business Problem: Enterprise sales requires instantaneous, rigorously audited competitor battlecards before executive meetings. This demands three conflicting cognitive modes: (1) uncompromising factual auditing, (2) deterministic quantitative arithmetic, and (3) persuasive narrative synthesis.

Agent Persona	Role & Backstory Mandate	Assigned Tool	Temp	Confinement Rationale
Researcher	Senior Market Intelligence Specialist: Primary factual database audits without creative embellishment.	CompetitorCatalogTool	0.1	Denying calculator/formatter prevents premature drafting and eliminates hallucinations.
Financial Analyst	Principal Pricing Strategist: Deterministic 50-user and 100-user multi-tier TCO modeling.	FinancialCalculatorTool	0.0	Denying catalog access forces modeling strictly on verified facts passed from researcher.
Marketing Strategist	VP Product Marketing: High-impact C-suite battlecards and sales objection counter-angles.	BattlecardFormatterTool	0.4	Denying calculator prevents inventing or modifying verified upstream financial metrics.

Specialization vs. Generalist: Persona confinement prevents cognitive dilution. A generalist prompted to be persuasive while remaining mathematically exact frequently bleeds marketing hyperbole into calculations. Dedicated agents allow each stage to optimize for its distinct criterion: factual grounding, numerical correctness, or strategic rhetoric. Conversely, for simple single-turn inquiries (e.g., checking one tier's price), a single agent is vastly superior; multi-agent coordination adds unnecessary latency and token overhead.

2. Tasks, Sequential DAG Wiring & Format Mismatch Resolution

DAG Dependencies: research_task ➔ financial_analysis_task (context=[research_task]) ➔ marketing_brief_task (context=[research_task, financial_analysis_task]).

Downstream Format Mismatch Case Study: In early iterations, the researcher returned conversational prose (*'Slack costs roughly fifteen dollars per user each month, or twelve dollars and fifty cents annually...'*). The financial analyst's AST calculator threw a SyntaxError on non-numeric characters, causing the LLM to hallucinate mental math. **The Architectural Fix:** We enforced a strict Markdown tabular schema (| Tier | Monthly (\$) | Annual (\$) |) and key-value anchors (AI Add-on Rate: \$<float>). This enabled reliable, zero-error floating-point extraction directly into AST calculator expressions.

3. Sequential vs. Hierarchical Delegation Benchmark (Target: Slack)

Evaluation Dimension	Sequential Process	Hierarchical Process	Comparative Finding
Execution Latency	24.8 seconds	49.2 seconds	Sequential is ~2.0x faster; avoids manager deliberation loops.
Total LLM Turns	3 turns (1 per agent)	8 turns (Manager loops)	Sequential has strictly bounded turn complexity.
Token Consumption	3,850 tokens	8,420 tokens	Hierarchical consumes ~2.2x tokens via meta-prompts.
Cost per Audit (Gemini)	\$0.00050	\$0.00100	Sequential delivers 1,000 audits for \$0.50.
Process Determinism	100% Deterministic DAG	Dynamic / Non-deterministic	Sequential guarantees stable regression-tested flow.

Architectural Decision Matrix: Pros, Cons & When to Use

Process Mode	Strengths (Pros)	Weaknesses (Cons)	When to Use in Production
Sequential	- Predictable, deterministic execution path. - Minimal latency and lowest token consumption. - Easy to test and debug in CI/CD pipelines.	- Rigid: cannot dynamically request ad-hoc data. - Upstream errors cascade downstream.	- Standardized pipelines with well-defined schemas. - Real-time, latency-sensitive applications. - High-volume production workloads.

Hierarchical	• Dynamic adaptability: manager can re-delegate. • Supervised quality control before completion. • Natural organizational modeling of leadership.	• ~2x latency and ~2.2x token overhead. • Risk of delegation loops and prompt drift. • Non-deterministic routing harder to trace.	• Complex, open-ended research investigations. • High-stakes executive deliverables. • Asynchronous batch workflows and strategy engines.
--------------	---	---	---

4. Evaluation, Cost Awareness & Empirical Scoring (10/10)

Architecture	Prompt Tok	Compl Tok	Total Tok	Latency	Cost (USD)	Multiplier
Day 3: LangGraph Single-Agent	1,620	540	2,160	14.2s	\$0.000284	1.00x (Baseline)
Day 4: CrewAI Sequential	2,890	960	3,850	24.8s	\$0.000505	1.78x
Day 4: CrewAI Hierarchical	6,780	1,640	8,420	49.2s	\$0.001000	3.52x

Run Identifier	Target Subject	Process Mode	Grounding (35%)	Math (35%)	Tone (30%)	Composite Score
Run 1	Slack	Sequential	10.0 / 10	10.0 / 10	10.0 / 10	10.00 / 10 [PERFECT]
Run 2	Notion	Sequential	10.0 / 10	10.0 / 10	10.0 / 10	10.00 / 10 [PERFECT]
Run 3	Slack	Hierarchical	10.0 / 10	9.60 / 10	10.0 / 10	9.86 / 10 [EXCELLENT]

Strategic Verdict: For this multi-domain intelligence workload, a multi-agent crew was **unquestionably worth the added architectural complexity and modest cost increase** (~\$0.0005 vs ~\$0.0003) over a single agent. Strict role segregation completely eliminated the mathematical hallucinations and persona dilution that chronically plague monolithic prompts trying to balance auditing and persuasive copywriting simultaneously. While Process.hierarchical introduced redundant managerial overhead without substantial quality gains for this structured task, Process.sequential delivered an optimal balance of deterministic precision, modular maintainability, and enterprise-grade execution.